

what is spark

spark is a distributed big data processing engine. it can process batch as well as real time

Spark Performance Tuning – Best Guidelines & Practices

- Use DataFrame/Dataset over RDD.
- Use coalesce() over repartition()
- Use mapPartitions() over map()
- Use Serialized data format's.
- Avoid UDF's (User Defined Functions)
- Caching data in memory.
- Reduce expensive Shuffle operations.
- Disable DEBUG & INFO Logging.

why dataframe is faster than rdd

RDD – RDD API is slower to perform simple grouping and aggregation operations. DataFrame – DataFrame API is very easy to use. It is faster for exploratory analysis, because rdd not

Spark – Difference between Cache and Persist?

Using cache() and persist() methods, Spark provides an optimization mechanism to store the intermediate

Both caching and persisting are used to save the Spark RDD, DataFrame, and Dataset's. But, the difference

PySpark Broadcast Variables

<https://sparkbyexamples.com/pyspark/pyspark-broadcast-variables/>

In PySpark RDD and DataFrame, Broadcast variables are read-only shared variables that are cached and available

follow->interview_preparations/spark_tutorial/Spark_RDD_Tutorial

diff between Accumulator and Broadcast Variables

An accumulator is also a variable that is broadcasted to the worker nodes.
The key difference between a broadcast variable and an accumulator is that while the broadcast variable is read-only, the accumulator can be added to

Accumulator can be used to implement counters (as in MapReduce) or sums.
Spark natively supports accumulators of numeric types, and programmers can add support for new types.

spark executor(imp)

<https://spark.apache.org/docs/latest/cluster-overview.html>

find second highest salary

```
select quantity from orders order by quantity desc limit 1 offset 1;
```

interview links:

apache spark interview question

<https://www.simplilearn.com/top-apache-spark-interview-questions-and-answers-article>

<https://sparkbyexamples.com/pyspark/pyspark-aggregate-functions/>

spark example run

<https://sparkbyexamples.com/spark/spark-how-to-kill-running-application/>

rdd of spark

<https://spark.apache.org/docs/latest/api/python/reference/api/pyspark.RDD.html>

increase and decrease no of partitions and parallelism

```
spark.conf.set("spark.sql.shuffle.partitions", "1000")
spark.conf.set("spark.default.parallelism", "1000")
```

When you run a PySpark RDD, DataFrame applications that have the Broadcast variables def

PySpark breaks the job into stages that have distributed shuffling and actions are executed. Later Stages are also broken into tasks.

Spark broadcasts the common data (reusable) needed by tasks within each stage.

The broadcasted data is cached in serialized format and deserialized before executing each task.

You should be creating and using broadcast variables for data that is shared across multiple tasks.

show full text in df

```
df.show(truncate=False)
```

spark coding links

<https://sparkbyexamples.com/spark/spark-web-ui-understanding/>

<https://spark.apache.org/docs/latest/api/python/reference/api/pyspark.RDD.html>

what is job scheduling in spark

By default, Spark's scheduler runs jobs in FIFO fashion. Each job is divided into "stage"

spark executor(imp) and spark architecture

<https://www.youtube.com/watch?v=JSvzGJC8vjQ>

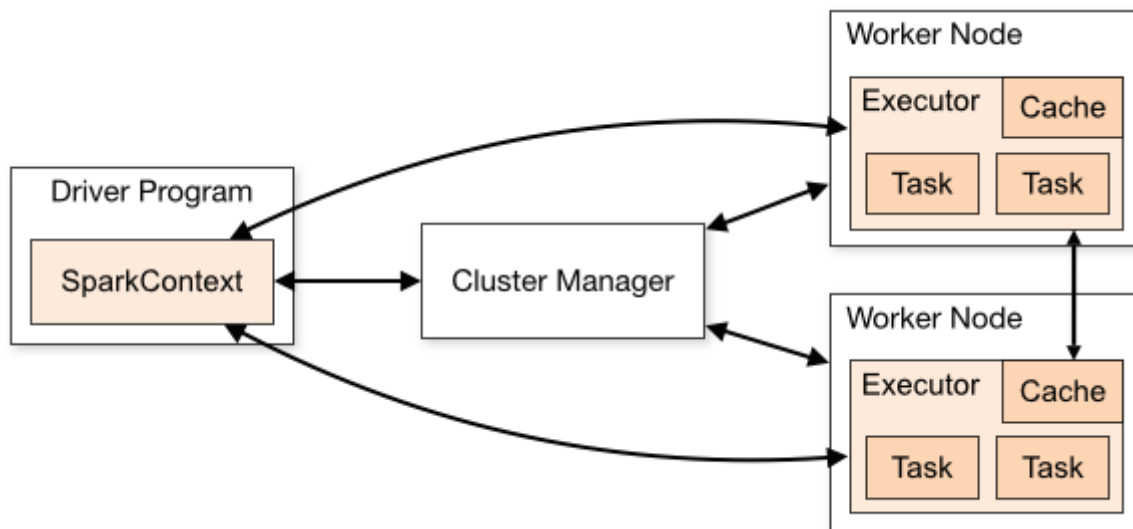
spark core component : spark ML-LIB, spark-GraphX, spark-streaming, spark-sql

spark executor : standalone(local), kubernetes, hadoop yarn, apache mesos

standalone : this is very simple cluster manager included with spark

hadoop-yarn : the resource manager of hadoop 2 which stands for yet another negotiator

apache-mesos : general cluster manager that can run hadoop mapreduce and services as well



spark Architecture define into three part

- 1.driver program/sparkSession
- 2.cluster manager
- 3.worker node (Executor + task)

1.driver program /spark session =>

this is the main entry point of spark application. the main method of program runs in driver
 driver create sparkSession/ sparkContext
 driver perform transformation and action

main role of driver program is that

1. convert user program into task
- 2.scheduling task in the executor

2.cluster manager =>

cluster manager are responsible for resource allocation and launch executor.
 cluster manager is pluggable component of spark
 cluster manager can dynamically adjust resource used by the spark application

3.Spark Executor:

the individual task given by the spark job are runs in the spark executor
 Executor Launch once in the beginning of the spark application and they run for the entire life cycle

the two main role play the executor

1. runs the task that makes up the spark application and return the result to the driver
2. provide in memory storage for RDD/DataFrame/DataSets that are cached by the users

how to optimize spark job

- 1.repartitions (repartitions split and shuffle the data across all the nodes in a cluster)
- 2.broadcasting
- 3.Avoid UDF and UDAF (please try to avoid user define function and use spark optimize functions)
- 4.Dynamic allocation

According to the workloads it is possible to scale up and scale down the number of executors

```
$ spark.dynamicAllocation.enabled //Set the value to true to enable dynamic allocation
```

5.Parallelism

In spark higher level APIs like dataframe and datasets use the following parameters

```
$ spark.sql.files.maxPartitionBytes //Specify maximum partition size by default 128MB
```

what is dataset in spark

Dataset is a new interface added in Spark 1.6 that provides the benefits of RDDs (strong typing, ability to use higher-level APIs like maps, etc.)

the basic difference between dataframe and dataset is

Type Safety: Dataset provides compile-time type safety. It means that the application's

. Python does not have the support for the Dataset API.

what is parquet and why it is used and how this is works internally

<https://www.upsolver.com/blog/apache-parquet-why-use>

what is row based and column based

let's take an example we have a file in which have 5 columns and 100 records then what happens

in csv make 100 files in which one full information of a row but other hand parquet store data into

pyspark join

pyspark dataframe support sql type join like inner, outer, left, right

```
df.join(df2, df.id==df2.id, 'inner')
```

mapPartitions()

is used to provide heavy initialization for each partition instead of applying to all elements this is the main difference

data is skewed to remove by using repartitions

Use option `maxRecordsPerFile` if you want to control the number of records for each partition. This is particularly helpful when your data is skewed (Having some partitions with very low records and other partitions with high number of records).

spark suffling and partitions

<https://sparkbyexamples.com/spark/spark-shuffle-partitions/>

Based on your dataset size, number of cores, and memory, Spark shuffling can benefit or harm your jobs. When you dealing with less amount of data, you should typically reduce the shuffle partitions otherwise you will end up with many partitioned files with a fewer number of records in each partition. which results in running many tasks with lesser data to process.

On other hand, when you have too much of data and having less number of partitions results in fewer longer running tasks and some times you may also get out of memory error.

Getting a right size of the shuffle partition is always tricky and takes many runs with different value to achieve the optimized number. This is one of the key property to look for when you have performance issues on Spark jobs.

diff between ETL AND ELT

diff.

Transformations are done in ETL server/staging area. but in ELT Transformations are performed in the target system

ETL model used for on-premises, relational and structured data.Used in scalable cloud infrastructure which supports structured, unstructured data sources.